

Noise-Trader Activity, Price Formation and Liquidity on the Tokyo Stock Exchange

A technical briefing on the prototype implementing the future work
proposed in Ohta (2026) — January–April 2025 tape

Kazumi Li

July 2026

Summary

This briefing documents a working prototype of the future work proposed at the end of Ohta (2026): treating price clustering as an observable daily proxy for noise-trader activity and relating it to price formation and liquidity. It is written for one reader — the author, preparing to present and extend this work during the MTEC internship — so it optimises for being explainable: Section 1 maps the document, Table 1 states every finding with its verdict, and Section 4 gives the mathematical foundation of each measure with the reference it comes from. The sample is the January–April 2025 Tokyo Stock Exchange tape: 77,592 stock-days, 1,046 stocks, 78 trading days, under the paper’s own sample rules.

The measurement layer replicates the paper and is the sturdiest part of the study. Round prices carry 14.95% of buy-initiated large-trade volume against 10.77% of small-trade volume (the paper: 14.8 and 11.1), the round-price impact premium is positive at about 34% of the paper’s published one-minute size, and the ladder-inferred submission shares land on the paper’s levels without calibration (0.102 vs 0.106) — while the cancellation ratios overshoot (0.91 vs 0.81), a difference reported as open rather than folded into the validation claim.

On the paper’s open question, clustering carries a small amount of information about next-day trading costs (0.0188 on the log effective spread, $t = 2.51$; depth moves the complementary way), but the pooled estimate is significant at the 5% level, but it falls below conventional significance once the outcome’s history means five lags rather than one; its sign reverses between the within-stock and cross-sectional designs. The within-day version is the robust one: bucket-level clustering predicts next-bucket spreads (0.0189, $t = 8.80$) under stock-day and bucket fixed effects. A placebo rebuilding the measure on each other digit leaves 0 of nine significant.

The strategy demonstration exists to price the idea, and it does: the daily long–short costs 32.0 basis points a day against a gross of -3.1. Sorting on a five-day mean of the signal instead cuts turnover from 74% to 31% of each leg per day and the cost line from 32.0 to 13.5 basis points, and the net line is still -15.5 – the churn is measurement noise, and even without it there is no alpha at this horizon. The economics point at execution scheduling — when and at which price levels to trade a rebalance that is happening anyway — not at standalone alpha.

Contents

1	What this briefing is, and how to read it	5
1.1	Purpose	5
1.2	The sentence this study is built on	5
1.3	Why the answer is not obvious in advance	5
1.4	What this study can and cannot do	6
1.5	The findings at a glance	6
1.6	Map of the document	6
2	Where each ingredient comes from	8
3	Data, institutions, and the sample	10
3.1	The feed	10
3.2	Two institutional details that shape everything	10
3.3	Building the sample	11
4	Mathematical foundations, and the validation they had to pass	12
4.1	Clustering (Ohta 2026)	13
4.2	Liquidity (the paper’s definitions)	13
4.3	Price formation	14
4.4	The visible ten-level book	14
4.5	Limit-order flow inferred from the ladder	14
4.6	Estimation and inference	15
4.7	What the measurement layer had to prove	15
5	Replication: does 2025 look like the paper’s sample?	16
5.1	The headline measures	16
5.2	The digit distribution	17
5.3	Finer grids, and the size question	17
5.4	The impact premium	18

5.5	Standing round-price depth	20
5.6	The contamination the paper warns about	22
6	Finding: clustering and next-day liquidity	22
6.1	Specification, and why it is the dynamic one	22
6.2	What the panel says	23
6.3	Which measure carries the information	23
7	Findings: the order book, and the day within the day	24
7.1	Rebuilding the paper's equations without its proxies	24
7.2	The shape of the trading day	26
7.3	Within the day	26
8	How fragile is the daily finding? Robustness	27
8.1	The placebo	27
8.2	Subsamples, deeper dynamics, and an alternative estimator	29
9	A costed strategy demonstration	30
10	What to take to MTEC	32
10.1	What was found	32
10.2	What it is not	33
10.3	Mapping the results onto the internship brief	33
10.4	Conclusion	34
	References	35
	Appendix	37
	Appendix A Variable definitions	37
	Appendix B Inferring limit-order flow from the visible ladder	40
B.1	The accounting identity	40
B.2	Two rules that make it correct	40
B.3	What this cannot see	41
B.4	The test	41
B.5	Does it recover the published magnitudes?	42
	Appendix C Reproducibility	43
C.1	Pipeline	43
C.2	The stop rule	43
C.3	Tests	43

C.4	Two decisions worth recording	44
C.5	Data	44

1 What this briefing is, and how to read it

1.1 Purpose

This document is a technical briefing, not a journal article. It has one reader — the author, preparing to present, defend and extend this work during an internship whose task is a rebalancing strategy built on the indicator studied here — and it is organised so that any result can be explained from it directly: what was computed, from which formula, taken from which reference, validated how, and trusted how far. Everything numerical in the chapters that follow is regenerated from the pipeline’s stored outputs whenever the analysis is re-run, and every interpretive sentence is generated from the estimates it describes, so the prose cannot drift from the tables.

1.2 The sentence this study is built on

Ohta (2026) asks what kind of participant behaviour produces *price clustering* — the disproportionate share of trading that executes at round prices — and answers it with an order-flow mechanism. Noise traders leave limit orders at round prices outside the best quote, because working out a precise valuation is expensive and a round number is a cheap substitute. They are slow to revise or cancel those orders, because managing orders at speed is also expensive. When value moves, faster participants take the stale orders with large market orders. Three things follow, and the paper documents all three: clustering concentrates in large trades rather than small ones, it is stronger where individual investors are more present, and trades at round prices carry about a basis point more price impact than trades elsewhere, because they execute against prices that have gone stale.

The paper then closes with a claim and a gap:

In light of these results, it may be possible to construct a measure of noise trader activity from price clustering measures. . . . In that case it may be possible to analyze the relationship between noise trader activity and price formation and liquidity at a daily or shorter frequency. Such analysis, however, remains a task for future work.

This study is that analysis, on a stretch of tape three years past the end of the paper’s sample.

1.3 Why the answer is not obvious in advance

In Kyle (1985) noise traders are what makes a market liquid: they are the volume informed traders hide inside, and without them the spread widens to protect against adverse selection. Peress and Schmidt (2020) find the empirical counterpart — when noise traders are distracted, liquidity gets worse. On that reading, high clustering should mark a *more* liquid market.

The mechanism Ohta documents points the other way. If clustering marks orders that are stale and about to be picked off, it marks a market where liquidity is being supplied by participants who are about to regret it, and where the price impact of a large order is unusually high. On that reading, high clustering should mark a market that is *more* expensive to trade.

Both stories are about the same participants. They differ on whether what matters is that the uninformed are present or that they are slow. Either answer would be a result, and the sections below report which way this sample falls — with the robustness work deciding how hard the answer may be leaned on.

1.4 What this study can and cannot do

Ohta identifies noise traders using two proxies for individual-investor activity: margin-trading balances and ownership shares. Neither is available here. That sounds like a handicap and is partly the point: the paper's own conclusion is that the clustering measures *are* the proxy, so a study that uses only clustering and the order book is testing the claim rather than assuming it.

Working without those proxies also forces a substitution that turns out to be an advantage. The data here carry the ten best price levels on each side of the book after every event. That makes it possible to see the standing round-price inventory directly — what share of visible resting depth is sitting at round prices at any moment — where a study confined to executed trades sees round-price orders only after someone has taken them. The same data support inferring limit-order submission and cancellation volume from how the displayed ladder changes, which is how Ohta builds his own order-flow variables.

The limits are equally plain. Four months is not a long panel. There is no experiment and no instrument, so nothing below is causal; the results are descriptive associations and predictive relationships, and are labelled as such throughout. The ten visible levels are a narrow window onto a book where most resting volume sits further out.

1.5 The findings at a glance

Table 1 states every finding with the verdict the robustness work assigns it. When presenting this study, that table is the honest summary; each row's section gives the design and the caveats in full.

1.6 Map of the document

Section 2 says where each ingredient of the study comes from — every borrowed idea with its source. Section 3 describes the feed, the institutional details that govern the measures, and how the sample is built, including when its filters bite. Section 4 is the mathematical foundation: every measure's definition, the reference it originates from, and the validation the implementation had to pass. Section 5 is the replication against the paper's published magnitudes. Section 6 is

Table 1: Findings and their verdicts, in one place

Finding	Headline estimate	Verdict
F1 Clustering levels replicate	$M^{BLarge0}$ 14.95% vs paper 14.8%	Matches; validates the measurement
F2 Round-price impact premium	+0.43/+0.51 bp at 60s	Positive, about 34% of the paper's one-minute size; gone by five minutes
F3 Size gradient	Q1 14.83% vs Q5 15.82% (pooled)	Absent in this screened universe
F4 Clustering → next-day spread	0.0188 ($t = 2.51$)	Fragile: it falls below conventional significance once the outcome's history means five lags rather than one; its sign reverses between the within-stock and cross-sectional designs
F5 Clustering → next-day depth	-0.0441 ($t = -5.22$)	Secondary outcome, same design and caveats
F6 Within-day (30-min buckets)	0.0189 ($t = 8.80$)	The robust version of the liquidity link
F7 Placebo on digits 1–9	0 of 9 significant	Result specific to the round digit
F8 Mechanism variables (L, L^C , RDepth)	RDepth $t = 11.4$; L, L^C null	1 of 3 sign predictions confirmed; flow shares censored at 10 levels
F9 Ladder inference vs paper	L^{S0} 0.102, L^{S0C} 0.91	Submission shares match (0.102 vs 0.106); cancellation ratios overshoot (0.91 vs 0.81)
F10 Strategy demonstration	net -35.1 bp/day daily; -15.5 smoothed	Costs dominate at daily horizon; smoothing cuts costs 32 to 14 bp/day, still negative net

Every verdict is generated from the estimates it describes; the sections give the designs and the caveats in full. Coefficients are in log points of the outcome per unit of the clustering share unless stated otherwise.

the daily clustering–liquidity relationship, Section 7 the order-book anatomy and the within-day results, and Section 8 the robustness work, including the placebo, the sub-period splits, and an estimator cross-check that disagrees informatively. Section 9 prices a naive strategy on the signal, in two variants. Section 10 maps everything onto the internship brief. The appendices carry the full variable definitions, the ladder-inference algorithm with its hand-built test, and the reproduction guide.

2 Where each ingredient comes from

Every borrowed idea in this study is listed here with its source, so that any formula or design choice in Section 4 can be traced to the paper it originates from in one hop.

Clustering as a measurement problem. Harris (1991) established the facts every later study works against: trades concentrate on round fractions, and that concentration rises with volatility and with the price level, falls with market capitalisation and trading frequency, and rises as the tick grid gets finer. The last of these matters directly here, because the Tokyo Stock Exchange operates two grids at once and moved a large part of its market onto the finer one in 2023. Chung et al. (2004) add the correlation with the bid–ask spread that makes spread a necessary control rather than an optional one.

Who does the clustering. The behavioural reading is that round numbers are cognitive shortcuts. Kuo et al. (2015) give it teeth by showing that the investors who use round-price limit orders most heavily also perform worst, which is what licenses reading a round-price order as an uninformed one. Bhattacharya et al. (2012) take the same observation in a tradable direction, documenting buy–sell imbalances immediately around round numbers. Davis et al. (2014) supply the fact that organises Ohta’s whole argument: clustering lives in *large* trades, not small ones.

Clustering in the book, not only in the tape. Ahn et al. (2005) look at the limit order book directly and find it thicker at round prices. That is the closest antecedent to the depth-based measure introduced in Section 4, and the difference is one of data: a study that sees only executed trades observes round-price orders after they have been taken, whereas a study that sees the standing book observes them while they are still exposed. Box and Griffith (2016) is the closest design sibling to the present paper, pairing buy/sell asymmetry in clustering with the price impact of round-price trades.

Are individuals noise traders at all? Ohta’s identification rests on treating individual-investor activity as a proxy for noise trading, and the evidence there is genuinely mixed. Barber and Odean (2008) find individuals buying on attention rather than information, which is the noise-trader

reading; Kaniel et al. (2008) find individual order flow predicting returns, which is not. This study sidesteps the question rather than settling it, because it uses no individual-activity proxy: it works from the clustering measures and the order book alone, which is what the paper's own conclusion proposes.

Yamada (2023) supplies the other half of the mechanism on this exchange specifically, identifying high-frequency traders on the TSE that submit market orders against round-price limit orders. That is the counterparty Ohta's account requires, observed directly.

Noise traders and liquidity. Whether noise trading helps or harms liquidity is genuinely unsettled, and this study cannot settle it. In Kyle (1985) noise traders are the cover under which informed traders operate, so more of them means a deeper market. Peress and Schmidt (2020) find the empirical counterpart: when noise traders are distracted, liquidity deteriorates. Against that, Foucault (1999) formalises pick-off risk — the cost of leaving a limit order unrevised — and Aquilina et al. (2022) measure the speed race that determines who bears it. If clustering marks orders that are about to be picked off, it may mark a market that is about to become more expensive to trade, not less. Section 6 reports which way this sample falls, and treats either answer as a result.

Reading price impact carefully. Ohta's second contribution is a warning about interpretation. Glosten and Milgrom (1985) and Hasbrouck (1991) make price impact the canonical measure of information asymmetry, and Hendershott et al. (2011) give the decomposition of the effective spread into impact and realized spread that this study uses. But in a limit order market, impact also rises when stale orders sit in the book waiting to be taken, which has nothing to do with anyone being informed. The impact premium at round prices reported in Section 5 is exactly that second channel.

Beyond the best quote. Biais et al. (1995) is the canonical anatomy of a pure limit order market of the kind the TSE runs. Cao et al. (2009) justify looking past the best quote at all, finding that the levels behind it contribute materially to price discovery; Næs and Skjeltorp (2006) turn the shape of those levels into the depth-slope measure used here. The order-flow imbalance construction is that of Cont et al. (2014).

Inference. Every panel regression below carries stock and day fixed effects with standard errors clustered on both margins, following Cameron et al. (2011). With the same stock appearing on nearly every day of the sample and the same day covering a thousand stocks, the alternative is not a slightly optimistic standard error but a badly wrong one. The dynamic specification inherits the small bias Nickell (1981) characterises, of order $1/T$; the estimator cross-check follows Fama and MacBeth (1973) with Newey and West (1987) standard errors over the daily coefficient

series, and is reported precisely because the two designs identify different comparisons — within a stock over time versus between stocks — so their agreement or disagreement is itself a finding.

3 Data, institutions, and the sample

3.1 The feed

The source is the Nikkei NEEDS TICST120 product for 2025: one record per market event, carrying the trade if there was one and the ten best price levels on each side of the book as they stood afterwards. Trade direction is supplied by the exchange rather than inferred, which removes a source of error that would otherwise sit underneath every result here. Daily summaries from the companion TICSS110 product supply trading units, shares outstanding, and the variables used to screen the universe.

The 2025 feed arrived complete: every trading day is present with a shard count in line with its month, and no date had to be excluded for a delivery gap. That gives 243 usable trading days in the year, of which this study uses the months listed below.

3.2 Two institutional details that shape everything

The trading session is not a fixed object. The afternoon session closed at 15:00 until 2024-11-05 and at 15:30 afterwards. This panel sits entirely after that change, so the close is 15:30 throughout — but the close is still resolved from the date rather than assumed, because the same code is meant to run on earlier years, and the price-impact horizons depend on knowing when the session actually ends.

The exchange runs two tick grids at once. Constituents of the TOPIX 500 trade on a finer grid than everything else — a regime extended from the TOPIX 100 to the whole index on 2023-06-05 and stable since. The schedule is encoded from the exchange’s published tables (Japan Exchange Group, 2024) and cross-checked against ticks read directly off the tape. The continuous-auction mechanism these measures are computed over is described in Lehmann and Modest (1994).

This interacts with Ohta’s sample filter in a way worth stating plainly, because it shapes the sample more than any other decision. The filter admits only tick sizes that are powers of ten, since the last-digit-of-ten construction is otherwise undefined — on a 0.5-yen grid the last digit takes only even values and “round” stops meaning what it means elsewhere. On the finer grid the 1,000–3,000 yen band is 0.5 yen, so a large part of the mid-priced TOPIX 500 population is excluded outright. The regression sample is therefore tilted away from mid-priced large caps, and the composition is reported rather than papered over:

Table 2: Tick-size composition of evaluated stock-days

Tick size	Stock-days	Status
1 yen	81,159	kept
10 yen	3,626	kept
0.1 yen	2,668	kept
0.5 yen	1,729	excluded by filter (d)
5 yen	1,260	excluded by filter (d)
3 yen	4	excluded by filter (d)

3.3 Building the sample

The universe is selected by Ohta's own criteria, applied first to the daily summaries so that the expensive tick-level work runs only on stocks that could enter the sample. He studies First Section (now Prime) common stocks whose trading days satisfy four conditions on more than half the year's sessions: the first trade by 9:10, an opening price above 200 yen, more than twenty continuous-session trades, and a tick size that is a power of ten throughout the day. Regional venues, exchange-traded funds and trust units — identifiable by their trading units — are excluded.

Each of those filters exists for a reason worth keeping in view. The 9:10 rule removes days that opened late after a price-limit halt, which would have a shortened continuous session and therefore a differently-behaved measure. The 200-yen floor removes stocks so cheap that the price cannot traverse all ten digits without moving several percent. The twenty-trade minimum removes days where the measure is a ratio of small integers.

Applying the same tests properly at tick level gives the final panel. Table 3 is the whole construction: no stock-day disappears without being counted, which is the only way a reader can tell selection from attrition.

Table 3: Sample construction

Step	Stock-days	% of attempted
Stock-days attempted	90,870	100.0
Survived the pre-screen	90,449	99.5
First trade by 9:10	89,319	98.3
Opening price above ¥200	90,429	99.5
More than 20 continuous-session trades	90,308	99.4
Tick a power of ten all day	81,363	89.5
Passing all four filters	80,254	88.3
Stock qualifies on more than half the sample	77,592	85.4

The four filters are Ohta's, applied at tick level. The pre-screen is a cheap four-column read that rejects stock-days which cannot possibly qualify — overwhelmingly those with too few trades — and changes nothing about the final sample. The filters are not nested, so their individual pass counts do not multiply to the joint count.

Coverage. The panel covers 78 trading days across 4 consecutive calendar months, from

2025-01-06 to 2025-04-30, rather than the full 243-day year. Ingesting and processing the complete tape for this universe takes substantially longer than the time available for a prototype, and a contiguous block was chosen over scattered months so that every lag in the panel is a genuine one-day lag.

That choice has a cost worth naming. Four consecutive months is one market regime, not a cross-section of them, and this particular block is not a quiet one. Sub-period stability is therefore tested directly, on halves derived from the sample itself, and the verdict lives with the evidence in Section 8 rather than being assumed here. Nothing about the pipeline changes with more months — the same command extends it and this document regenerates from whatever the panel contains — but the results below describe this stretch of 2025, not the year.

When the filters bite. Ohta’s 9:10 first-trade rule removes days that opened late after a price-limit halt. On an ordinary day that is a handful of names; on a market-wide shock it is a large slice of the universe at once. The thinnest days of this sample make the point: 627 of 1,165 candidate stocks survive on 2025-04-07 and 727 of 1,165 on 2025-04-10, against 86% on the median day. The panel therefore thins exactly when noise-trader activity was most extreme, and every daily estimate below is conditional on a day that opened normally. This is the paper’s own admission rule, inherited deliberately – but over four months containing such an episode it shapes the sample, and a reader comparing April to January should know it.

Verdict. The sample is Ohta’s, reconstructed on a year he did not have. Its two distinctive features — the exclusion of the 0.5-yen grid, and the thinning on shock days — both follow from his filters rather than from choices made here, and both are carried explicitly into the interpretation of every result below.

4 Mathematical foundations, and the validation they had to pass

This section is the reference core of the briefing: every quantity the study computes, as a formula, with the source it originates from and the implementation detail that makes it correct on this exchange. The full variable list is in Appendix A; the ladder algorithm has its own appendix. Throughout, i indexes stocks, t trading days, and prices are handled internally as integers in tenths of a yen, so the modular arithmetic below is exact on every grid — floating-point mod silently corrupts the 0.1-yen grid, which is exactly where clustering is strongest.

4.1 Clustering (Ohta 2026)

For tick size τ , the *last price digit* of a trade at price p is

$$d(p; \tau) = \frac{p \bmod 10\tau}{\tau} \in \{0, 1, \dots, 9\}, \quad (1)$$

and a *round price* is one with $d(p; \tau) = 0$ (Ohta, 2026). The construction requires the tick to be a power of ten — the paper’s sample filter (d) — and under a uniform digit distribution each digit carries 10% of volume, which is the benchmark every clustering number in this study is read against. The stock-day measures are volume-weighted shares at the round price,

$$M_{i,t}^0 = \frac{\sum_{\text{trades}} v \mathbb{I}[d(p; \tau) = 0]}{\sum_{\text{trades}} v}, \quad M_{i,t}^{k \text{ Size } 0} = \frac{\sum_{k, \text{Size}} v \mathbb{I}[d(p; \tau) = 0]}{\sum_{k, \text{Size}} v}, \quad (2)$$

with initiator $k \in \{B, S\}$ from the exchange’s own execution-type field (so no Lee and Ready (1991) inference is needed), *Large* meaning more than one trading unit and *Small* exactly one, following the paper. Only continuous-session (zaraba) trades executed against a two-sided ordinary quote enter; auction prints are excluded because their prices are not set by individual limit orders. The determinants of clustering — volatility, price level, capitalisation, trading frequency, tick — come from Harris (1991), which is where the control set originates.

4.2 Liquidity (the paper’s definitions)

With m the midquote prevailing immediately before a trade and $q = +1$ (-1) for buy-initiated (sell-initiated) trades, the *effective half-spread*, *price impact* at horizon Δ , and *realized spread* are the volume-weighted means of

$$es = q \frac{p - m_t}{m_t}, \quad imp_{\Delta} = q \frac{m_{t+\Delta} - m_t}{m_t}, \quad rs_{\Delta} = es - imp_{\Delta}, \quad (3)$$

in basis points, at $\Delta \in \{1s, 60s, 300s\}$. The decomposition of the effective spread into impact plus realized spread follows Hendershott et al. (2011); reading impact as the informational content of a trade goes back to Glosten and Milgrom (1985) and Hasbrouck (1991), and Ohta’s point — adopted here — is that in a limit-order market impact also rises when stale orders sit waiting to be taken, which is not information. Note there is no factor of two: these are half-spreads, matching the paper’s convention. A trade’s horizon must end inside its own session (no impact is measured across the lunch break or the close), and $RS = ES - Imp$ holds as an identity on each horizon’s own trade set — verified to three decimals as an implementation check, not a finding. The round-price impact premium is

$$\Delta Imp_{i,t}^{k \text{ Large}} = Imp_{i,t}^{k \text{ Large}, d=0} - Imp_{i,t}^{k \text{ Large}, d \neq 0}, \quad (4)$$

with cells of fewer than five trades set missing rather than reported. Box and Griffith (2016) is the closest antecedent for pairing clustering asymmetries with round-price impact. The auxiliary Amihud measure is $|r_{i,t}^{oc}|$ over turnover in millions of yen (Amihud, 2002).

4.3 Price formation

Realised variance is the sum of squared five-minute midquote log returns within each session, and the variance-ratio deviation is

$$|VR_{i,t} - 1|, \quad VR_{i,t} = \frac{\sigma_{5\min}^2}{5 \sigma_{1\min}^2}, \quad (5)$$

the standard variance-ratio construction measuring departure from a random walk in either direction; returns never span the lunch break or the overnight gap.

4.4 The visible ten-level book

Two measures use data the paper did not have. The *round-price resting depth* is the share of visible ten-level depth standing at round prices,

$$RDepth_{i,t}^{\text{ask}0} = \text{time-avg} \frac{\sum_{\ell \leq 10} v_{\ell}^{\text{ask}} \mathbb{I}[d(p_{\ell}^{\text{ask}}; \tau) = 0]}{\sum_{\ell \leq 10} v_{\ell}^{\text{ask}}}, \quad (6)$$

sampled on a one-minute grid over ordinary two-sided snapshots (uniform benchmark again 10%). The idea that the book itself is thicker at round prices is documented for Hong Kong by Ahn et al. (2005); the measure here is that observation turned into a stock-day share. The *order-flow imbalance* at the best quote follows Cont et al. (2014) exactly, and the *depth slope* — how fast size grows walking away from the best quote — follows Næs and Skjeltorp (2006); Cao et al. (2009) is the license for looking past the best quote at all.

4.5 Limit-order flow inferred from the ladder

Ohta's explanatory variables are the round-price shares of submitted and cancelled limit-order volume. The feed does not carry orders, so they are inferred from the accounting identity that resting volume at a price can only change by execution, cancellation or submission: with $\delta = V_n(p) - V_{n-1}(p)$ between consecutive snapshots and $X_n(p)$ the volume printed at p in the interval,

$$\text{submitted}_n(p) = \max\{\delta + X_n(p), 0\}, \quad \text{cancelled}_n(p) = \max\{-(\delta + X_n(p)), 0\}. \quad (7)$$

L^{k0} (round-price share of submissions) and L^{k0C} (cancellations over submissions at round prices) aggregate these to the stock-day, targeting the same quantities the paper builds from its feed.

Two rules make equation (9) correct in practice — match by price, never by level index, and count a price only where it is observable in both snapshots — and Appendix B states them, the three known failure modes, and the hand-built order book that tests all three.

4.6 Estimation and inference

The daily specifications are two-way fixed-effects panels,

$$y_{i,t} = \beta M_{i,t-1}^{BLarge0} + \gamma y_{i,t-1} + \theta' \mathbf{x}_{i,t-1} + \alpha_i + \delta_t + \varepsilon_{i,t}, \quad (8)$$

with standard errors clustered by stock and by day following Cameron et al. (2011) — with the same stock appearing on nearly every day and the same day covering a thousand stocks, one-way clustering is not conservative, it is wrong. The control vector $\mathbf{x}$ holds the relative tick, lagged log turnover, lagged log realised variance, lagged log spread, the overnight return, and the opening-digit contamination interactions of Section 5, timed to match the regressor. Including the lagged outcome makes (8) a dynamic panel; the induced bias is of order $1/T$ (Nickell, 1981) and negligible at T near eighty, and the robustness section re-estimates with five lags so that “the outcome’s own history” does not secretly mean one day. The estimator cross-check runs the same regression cross-sectionally day by day and averages the coefficients (Fama & MacBeth, 1973) with Newey and West (1987) standard errors over the daily series; agreement or disagreement between (8) and that average is informative about whether a relationship lives within stocks over time or between stocks, and Section 8 reports which it is here. Continuous variables are winsorised within stock at the 2.5th and 97.5th percentiles, the paper’s convention; the intraday specifications replace $\alpha_i + \delta_t$ with stock-day and bucket-of-day effects and cluster by stock and by date.

4.7 What the measurement layer had to prove

The measure library is a single module, and the full panel build calls it unchanged; validating it is validating production code. It had to pass four things.

Two hand-verified stock-days, to the second decimal. Two stock-days were computed independently in advance and used as tripwires: one on the 1-yen grid, one on the 0.1-yen grid where the digit arithmetic is fragile. All ten clustering cells reproduce to within 0.005 percentage points, and both tick sizes were recovered identically by the coded exchange schedule and by inference from the tape.

A negative control. Excluding the opening and closing auctions is not cosmetic. Letting them back in moves the daily measure by -2.91 and $+8.73$ percentage points on the two anchor days: the closing call is a single enormous volume-weighted print, and its price is not set by an

individual limit order the way continuous-session prices are. A pipeline that silently included it would produce measures that looked plausible and were wrong.

A hand-built order book. The ladder inference is checked against a six-snapshot tape whose order flow is known exactly, constructed so that each of the three standard failure modes produces a visibly wrong answer — mistaking an execution for a cancellation, inventing flow when the ladder shifts a level, and counting depth that entered from beyond the visible window. Appendix B gives the construction.

Magnitudes nobody tuned for — reported with their one miss. The strongest check is external. Ohta reports a mean round-price share of submitted sell limit orders of 0.106 and a mean cancellation ratio of 0.806 over 2010–2022. Run on this study’s panel, the inference lands on the submission share to half a percentage point without calibration — and overshoots the cancellation ratio (Appendix B carries the numbers). The agreement on levels is treated as validation; the overshoot on the ratio is treated as an open difference, not absorbed into the claim.

Verdict. The measurement layer reproduces two independently computed stock-days exactly, satisfies the identities it should satisfy, orders itself correctly across the liquidity spectrum, and lands on the published submission-share magnitudes it was never shown. The results that follow rest on it.

5 Replication: does 2025 look like the paper’s sample?

Nothing below this point is worth reading if the measurement is wrong, and the cheapest way to find out is to check whether a year the paper never saw reproduces the facts it established. This section is that check. It is also the first out-of-sample evidence on these measures.

5.1 The headline measures

Table 4 puts the 2025 means beside Ohta’s 2010–2022 benchmarks. Round prices take 14.22% of all continuous-session volume against his 13.5%, and 14.95% of buy-initiated large-trade volume against his 14.8%. Under a uniform digit distribution every one of these would be 10%. Clustering is alive and well on the Tokyo Stock Exchange three years past the end of his sample.

The ordering the mechanism predicts. The paper’s account is specifically about *large* orders taking stale round-price limit orders, so the large-trade measure should exceed the small-trade one. It does, by 4.18 percentage points on the buy side, significant at the 1% level. The small-trade measures sit at 10.77% and 10.95%, barely above the uniform benchmark. Whatever produces clustering, it is not something that happens to trades of one unit.

Table 4: Price clustering on the Tokyo Stock Exchange in 2025, against the published 2010–2022 benchmarks

Measure	Sample mean (%)	SE	Ohta (%)	Difference	Stock-days
M^0	14.22	(0.104)	13.5	+0.72	77,592
$M^{BLarge0}$	14.95	(0.179)	14.8	+0.15	77,407
$M^{SLarge0}$	15.83	(0.235)	14.3	+1.53	77,464
$M^{BSmall0}$	10.77	(0.051)	11.1	-0.33	77,592
$M^{SSmall0}$	10.95	(0.058)	10.9	+0.05	77,592

Volume-weighted share of trading executing at a last price digit of zero, by trade initiator and trade size. Large means a trade larger than one trading unit. Standard errors are clustered by stock and by day. The benchmark column reports Ohta’s (2026) 2010–2022 means. Under a uniform digit distribution every measure would be 10%.

5.2 The digit distribution




Figure 1: Volume-weighted share of trading at each last price digit, by tick regime. The dashed line is the 10% a uniform distribution would give.

Figure 1 is the whole phenomenon in one picture. Digit zero stands well above the uniform line; digit five stands slightly above it, which is what one expects if half-way points attract some of the same behaviour; and the remaining seven digits sit flat. That shape is itself a diagnostic — a measurement error in the digit arithmetic would not produce a clean spike at zero and a smaller one at five, it would produce something lumpy.

5.3 Finer grids, and the size question

Harris (1991) predicts more clustering on finer grids, and the 0.1-yen days deliver it: 18.99% against 14.15% on the 1-yen grid. The mechanism is not mysterious. At a price around 900 yen a 0.1-yen tick is about one basis point, fine enough that order placement collapses back onto

Table 5: Price clustering by tick size

Tick (¥)	M^0 (%)	$M^{BLarge0}$ (%)	SE	Stock-days	Stocks
0.1	18.99	19.05	(0.611)	2,095	30
1	14.15	14.87	(0.180)	72,489	978
10	12.45	13.87	(0.486)	2,823	43

Ohta's sample filter admits only tick sizes that are powers of ten, since the last-digit construction is otherwise undefined. Harris (1991) predicts more clustering on finer grids.

Table 6: Price clustering by market-capitalisation quintile

Quintile	M^0 (%)	$M^{BLarge0}$ (%)	SE	$M^{BLarge0}$, ¥1 grid (%)	0.1-tick share (%)	Stock-days
Q1	14.45	14.83	(0.304)	14.83	0.0	15,743
Q2	14.66	15.55	(0.291)	15.55	0.0	15,689
Q3	14.20	14.77	(0.257)	14.77	0.0	15,683
Q4	13.12	13.81	(0.220)	13.70	1.9	15,450
Q5	14.70	15.82	(0.250)	15.51	12.2	14,764

Quintiles are formed once per stock on its median market capitalisation over the sample, so a stock does not drift between quintiles as its price moves. Q1 is the smallest. The pooled columns mix tick regimes; the ¥1-grid column removes the mechanical elevation that 0.1-yen days carry, and the 0.1-tick share column shows how unevenly those days fall across quintiles.

whole-yen prices — and a whole yen is the roundest number available. These days are kept here, where they are a result, and excluded from the regressions, where the paper excludes them.

Table 6 is the one place this sample departs from the paper. Ohta reports clustering falling with market capitalisation; here the smallest quintile shows 14.83% and the largest 15.82%, with no reliable gradient between them. Part of Q5's pooled elevation is composition rather than behaviour: 12.2% of its stock-days sit on the 0.1-yen grid, whose clustering is mechanically higher, and restricting every quintile to the 1-yen grid moves Q5 from 15.82% to 15.51% against Q1's 14.83%. A genuine dip at Q4 (13.81%) remains that neither explanation predicts.

Two candidate explanations survive that decomposition, and neither is tested here. The screen keeps only stocks liquid enough to clear the paper's twenty-trade and tick filters on more than half the sample, which removes the small-cap end where his gradient is steepest; and four months of cross-section is thin for sorting on a slow characteristic. Settling it would take the Standard and Growth sections, which this study does not cover. The gradient is treated as absent in this universe, not as explained.

It does not change the strategy design in Section 9: size is neutralised there regardless, because the concern is that a sort might load on size, not that it must.


5.4 The impact premium

This is the paper's second result and the one that makes clustering matter for execution. A trade at a round price is executing against an order whose price has gone stale, so the midquote should

Table 7: The round-price price-impact premium in large trades

Horizon	Side	Pooled	SE	t	Stock-day mean	Stock-days	Ohta
1s	buy	0.737***	(0.027)	27.10	0.566	47,832	–
1s	sell	0.814***	(0.026)	31.11	0.714	50,469	–
60s	buy	0.428***	(0.052)	8.28	0.367	47,775	1.26
60s	sell	0.507***	(0.049)	10.34	0.606	50,423	1.12
300s	buy	-0.048	(0.089)	-0.53	-0.170	47,601	–
300s	sell	0.216**	(0.091)	2.37	0.310	50,258	–

ΔImp is the volume-weighted midquote change after trades at a last price digit of zero minus the same quantity for other trades, within large trades on each side, in basis points and signed so that a positive value means the price moved the way the initiator pushed it. The pooled column removes stock and day effects and weights each stock-day by the smaller of its two cell counts; standard errors are clustered by stock and by day. Cells with fewer than five trades are dropped. The final column is Ohta's (2026) 2010–2022 mean at the one-minute horizon. * $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$.

**Figure 2:** Round-price price-impact premium in large trades, with 95% confidence intervals from standard errors clustered by stock and by day.

move further afterwards. At the one-minute horizon the premium is 0.43 basis points on the buy side and 0.51 on the sell side — positive and precisely estimated, at about 34% of the 1.26 and 1.12 the paper reports for 2010–2022. A smaller premium in a faster market is what stale-order pick-off would predict, but that reading is not tested here; the honest statement is that the sign replicates and the magnitude does not. The premium is front-loaded: 0.74 basis points at one second, 0.43 at one minute, and no longer distinguishable from zero on the buy side at five minutes (-0.05). Whatever a round-price large trade pays extra, the midquote gives most of it back within minutes — transient pressure, not information. That shape matters for any execution use of the result: the cost of crossing at a round price is real but short-lived.

The effective-spread decomposition closes as an identity, which is a check on the implementation rather than a finding: impact plus realized spread equals the effective spread to three decimal places at every horizon tested.

5.5 Standing round-price depth




Figure 3: Binned means of the round-price share of visible ask depth against measured clustering in buy-initiated large trades.

The measure the paper’s data could not support says the same thing from the other side. 13.21% of visible ten-level ask depth sits at round prices against a uniform benchmark of 10%, and 14.00% on the bid. Figure 3 relates it to the executed measure: stock-days where more of the standing book sits at round prices are stock-days where more round-price volume executes. Where the clustering measures observe the mechanism after the fact, this observes the inventory while it is still exposed.

Table 8: Opening-price digit, volatility, and measured clustering

Opening digit	Low volatility (%)	High volatility (%)	Difference	Stock-days, low vol.
0	15.32	15.40	-0.07	7,760
1	14.45	14.56	-0.11	3,620
2	13.95	14.15	-0.20	3,207
3	13.71	14.09	-0.37	3,175
4	13.65	13.86	-0.21	3,093
5	13.58	13.83	-0.25	4,644
6	13.48	13.84	-0.36	3,212
7	13.22	13.83	-0.61	3,154
8	13.90	14.33	-0.42	3,589
9	14.19	14.27	-0.07	3,323

Mean M^0 by the last digit of the opening price, split on whether the *same day*'s realised variance was below that day's cross-sectional median. On quiet days a stock that opens far from a round price may never reach one, which depresses the measure mechanically. A median split is a diluted version of the paper's stuck-price notion – the median day still moves through many grid points – so Table 9 repeats the comparison on the bottom decile of the day's price range, where the mechanism has to show if it is real. The regressions control for the opening-digit by low-volatility interactions with the timing matched to the regressor.

Table 9: The contamination on genuinely stuck days

Opening digit	Stuck days (%)	Mobile days (%)	Difference	Stock-days, stuck
0	16.44	15.24	+1.20	1,532
1	15.47	14.39	+1.07	738
2	14.08	14.05	+0.03	660
3	13.24	13.98	-0.74	661
4	12.80	13.87	-1.07	665
5	12.78	13.80	-1.02	935
6	11.89	13.87	-1.97	661
7	11.69	13.73	-2.04	624
8	13.78	14.15	-0.37	735
9	14.98	14.15	+0.83	628

Mean M^0 by the last digit of the opening price. Stuck means the day's high–low range, scaled by the opening price, falls in the bottom decile of that day's cross-section – the closest available analogue of Ohta's Table 2 conditioning. The signature of the mechanical contamination is a positive difference at digit zero and negative differences in the middle digits: a stuck day prints mostly its opening neighbourhood, whatever digit that happens to be.

5.6 The contamination the paper warns about

Ohta’s Table 2 documents a mechanical distortion: a day that opens far from a round price and then barely moves can never print one, so its measure is low for reasons that have nothing to do with noise traders. Both ingredients are visible here. Opening prices cluster on their own — digit zero takes about a fifth of openings — and days opening within one digit of a round price measure 14.85% against 13.57% for mid-digit openers on low-volatility days, with a nearly identical gap (14.93% vs 13.84%) on volatile ones: under a median volatility split the level effect dominates and the interaction is faint, because the median day still traverses many grid points. On genuinely stuck days — the bottom decile of the daily price range, the closest analogue of the paper’s conditioning — the signature sharpens the way Table 2 says it should: a stuck day that opened on a round price measures 16.44% against 15.24% when mobile, while stuck days opening at digits four to seven measure 12.29% against 13.82% (Table 9). Every regression below therefore carries the opening-digit by low-volatility interactions with the timing matched to the regressor, and the strategy signal is residualised on the same-day set.

Verdict. The sample reproduces the paper’s core facts — clustering well above the uniform benchmark, concentrated in large trades, stronger on finer grids, without the paper’s size gradient in this screened universe, carrying a positive round-price impact premium. The one departure (the size gradient) is documented with its decomposition rather than smoothed over. The measurement is sound and the rest of the study can be built on it.

6 Finding: clustering and next-day liquidity

This is the question the paper leaves open. Clustering measures noise-trader activity; what does that activity do to the cost of trading?

6.1 Specification, and why it is the dynamic one

Three specifications are estimated for each outcome and Table 10 reports two of them. All carry stock and day fixed effects, the full control set, and standard errors clustered on both margins.

The *predictive* column regresses today’s liquidity on yesterday’s clustering. It fixes the ordering but not much else: both series are persistent, so a stock that is illiquid and clustered today was probably both yesterday, and the coefficient partly recovers that.

The *dynamic* column adds the outcome’s own lag, and it is the one the argument rests on. Its coefficient answers a sharper question — does yesterday’s clustering say anything about today’s liquidity that the outcome’s *own one-day lag* did not already say? With 78 time periods the resulting dynamic-panel bias is of order $1/T$ (Nickell, 1981) and can be neglected. One lag is not the same as “the outcome’s history”, so Section 8 re-asks the question against five lags.

Table 10: Price clustering and next-day liquidity

Outcome	Predictive	SE	Dynamic	SE	Stock-days
Effective spread (log)	0.0367***	(0.0080)	0.0188**	(0.0075)	74,034
Price impact, 60s (bp)	0.1587	(0.1476)	0.0417	(0.1498)	74,034
Realized spread, 60s (bp) [†]	-0.0039	(0.1705)	0.0776	(0.1691)	74,034
Depth at best quote (log) [†]	-0.0616***	(0.0092)	-0.0441***	(0.0085)	74,034
Realised variance (log) [†]	-0.0114	(0.0187)	-0.0065	(0.0177)	74,030
Variance-ratio deviation [†]	0.0106**	(0.0050)	0.0106**	(0.0050)	74,032
Amihud illiquidity (log) [†]	0.0372	(0.0301)	0.0305	(0.0317)	69,491

Each cell is the coefficient on $M_{t-1}^{B\text{Large}0}$, the round-price share of buy-initiated large-trade volume. All specifications carry stock and day fixed effects and the full control set (relative tick size, lagged log turnover, lagged log realised variance, lagged log effective spread, the overnight return, and the opening-digit by low-volatility interactions, timed to match the regressor: lagged interactions beside M_{t-1} , same-day ones beside M_t). The dynamic column adds the outcome's own lag, so its coefficient is the incremental predictive content of clustering given the outcome's own one-day lag; Section 8 probes deeper lags. Standard errors are clustered by stock and by day. Days on the 0.1-yen grid are excluded, following Ohta (2026). [†] secondary outcome, not pre-specified. * $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$. These are predictive associations in a single quarter, not causal estimates.

6.2 What the panel says

The coefficient on lagged clustering in the effective-spread regression is 0.0188, significant at the 5% level. Read economically, a one-standard-deviation increase in the round-price share of large-trade volume is associated with an effective spread about 0.2% higher the following day. For the price-impact outcome the coefficient is 0.0417, not statistically distinguishable from zero. Depth at the best quote moves the complementary way (-0.0441, significant at the 1% level): the day after heavy round-price trading quotes are wider and thinner.

Which way the open question falls. Section 1 set out two incompatible predictions. On the Kyle (1985) reading, more noise trading means a deeper market and cheaper trading. On the stale-order reading, clustering marks liquidity supplied by participants about to be picked off, and trading should get more expensive. The sign here is positive, favouring the second reading: days following heavy round-price trading are days when trading costs more. That is one sample of one market, and it is reported as an association rather than a mechanism — and where the coefficient is not distinguishable from zero, as a direction the data do not resolve.

6.3 Which measure carries the information

Table 11 substitutes the other clustering measures. If the relationship works through the channel Ohta describes — large orders taking stale round-price limit orders — it should live in the large-trade measures and not in the small-trade ones. This is a weaker test than the placebo in Section 8, but it uses the paper's own internal contrast, and it deserves an explicit reading rather than a table alone.

Table 11: Which clustering measure carries the information?

Regressor (lagged)	Coefficient	SE	Stock-days
$M^{BLarge0}$	0.0188**	(0.0075)	74,034
$M^{SLarge0}$	0.0255***	(0.0086)	74,090
$M^{BSmall0}$	0.0211**	(0.0093)	74,218
$M^{SSmall0}$	0.0158	(0.0108)	74,218
M^0	0.0418***	(0.0123)	74,218

Outcome is the log effective spread; the dynamic specification of Table 10 with the clustering measure replaced. The mechanism concerns large trades executing against stale round-price limit orders, so the large-trade measures are where an effect should appear.

The contrast the mechanism predicts does not appear. Per standard deviation of the sorting measure, the small-trade share moves next-day spreads by 0.16% against the large-trade share’s 0.23% — comparable, not smaller — and the pooled measure M^0 carries the largest coefficient of all (0.0418, $t = 3.41$). Elevated round-price trading of any kind predicts wider spreads; the specific stale-limit-order channel, which lives in large trades, is not what separates the measures here. That reads as a caution against the mechanism interpretation of the daily association, not against the association itself.

Verdict. Clustering carries a measurable association with next-day liquidity beyond the outcome’s own one-day lag, in a specification that absorbs every stock-level and market-wide effect. Section 8 bounds how hard this can be leaned on: the pooled estimate is significant at the 5% level, but it falls below conventional significance once the outcome’s history means five lags rather than one; its sign reverses between the within-stock and cross-sectional designs. The association is descriptive: no experiment, and no claim that intervening on clustering would move spreads.

7 Findings: the order book, and the day within the day

7.1 Rebuilding the paper’s equations without its proxies

Ohta explains clustering with limit-order submissions, their cancellation rate, and two proxies for individual activity. The proxies are unavailable here, so Table 12 substitutes what the book itself shows: the round-price share of submitted limit-order volume, the ratio of cancellations to submissions at round prices, and the round-price share of standing visible depth.

The hypothesis makes three sign predictions, and one of the three is confirmed at the 5% level. More round-price submissions should mean more round-price trading, so the coefficient on L^{S0} should be positive: it is 0.0041, not statistically distinguishable from zero. Orders that are *not* cancelled are the ones that get picked off, so the coefficient on the cancellation ratio

Table 12: Order-book anatomy of price clustering

Dependent variable	L^0	SE	L^{0C}	SE	RDepth	SE	Stock-days
$M^{BLarge0}$	0.0041	(0.0530)	-0.0098	(0.0137)	0.2557***	(0.0224)	9,527
ΔImp^{BLarge}	3.7891	(9.2906)	-0.6881	(2.0254)	17.6160*	(9.9998)	5,366
$M^{SLarge0}$	0.1479	(0.0969)	0.0118	(0.0198)	0.1766***	(0.0405)	9,537
ΔImp^{SLarge}	6.0420	(7.7278)	-2.4246	(2.1302)	4.1300	(4.0221)	5,722

Analogues of Ohta's (2026) equations (1) and (2), with book-derived variables replacing the margin-trading and ownership proxies, which are not available here. L^0 is the round-price share of limit-order volume submitted within the visible book, L^{0C} the ratio of cancelled to submitted volume at round prices, and RDepth the round-price share of visible ten-level resting depth. Sell-side book variables are paired with buy-initiated outcomes, since a buy market order executes against sell limit orders. All regressors are lagged one day. Stock and day fixed effects, standard errors clustered by stock and by day. The hypothesis predicts a positive coefficient on L^0 and RDepth and a negative one on L^{0C} . * $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$.

should be negative: it is -0.0098, not statistically distinguishable from zero. And more standing round-price inventory should mean more round-price execution: the coefficient on RDepth is 0.2557, significant at the 1% level.

The two flow-based nulls need a caveat before they are read as evidence against the mechanism. The inference watches ten visible levels while 81% of resting volume sits beyond them (Appendix B), so the daily L and L^C series carry heavy measurement error, and classical measurement error drives coefficients toward zero. A null here is what censoring would produce even if the mechanism held; these estimates do not adjudicate it. The depth variable below requires no inference at all, which is why it is the cleaner test.

Table 13: Standing round-price depth and next-day liquidity

Outcome	RDepth ask	SE	RDepth bid	SE	Stock-days
Effective spread (log)	0.1314***	(0.0169)	0.0433**	(0.0208)	74,218
Price impact, 60s (bp)	1.1252***	(0.4207)	-0.3255	(0.4168)	74,218
Depth at best (log)	-0.0153	(0.0231)	0.0522*	(0.0290)	74,218

RDepth is the share of visible ten-level resting volume sitting at round prices, time-averaged over the day and lagged one day. Dynamic specification, controls and fixed effects as in Table 10.

Table 13 asks the more direct question — whether standing round-price depth predicts liquidity by itself. This is the measure with the cleanest interpretation in the whole study, because it requires no inference at all: it is a share of displayed volume, read off the book.

One reading discipline before celebrating that coefficient: standing depth predicting execution at the same prices is partly mechanical — trades print where volume rests, whoever put it there and for whatever reason. The result is consistent with the stale-order account rather than diagnostic of it.

7.2 The shape of the trading day

Before turning to the regressions, the raw intraday profile is worth reporting on its own, because it reproduces a result from a different paper by the same author. Ohta (2006) finds that clustering on this exchange is highest at the open and decays through the session, which is what the price-resolution hypothesis predicts: uncertainty about value is greatest when trading starts, and a round number is the cheapest available substitute for an estimate.

Figure 4 shows the same curve two decades later. The round-price share opens at 16.4% and falls to 11.5% in the closing bucket, a decline of 4.9 percentage points, and it does so nearly monotonically. The effective spread traces the familiar intraday pattern beside it. Neither series was tuned to produce this, and the agreement with a 2006 result computed from a different sample is a further sign that the measurement is reading the phenomenon rather than an artefact.

It also sharpens the interpretation of what follows. If clustering were simply a proxy for wide spreads, the two curves would be redundant; they are not, and the regressions below control for the spread directly.




Figure 4: Round-price share of traded volume and the effective half-spread, by thirty-minute bucket. The dashed horizontal line is the 10% a uniform digit distribution would give.

7.3 Within the day

The paper asks for the relationship “at a daily or shorter frequency”. Table 14 is the shorter one: thirty-minute buckets, with stock-day fixed effects absorbing everything about the day — the stock, the news, the market state — and bucket-of-day fixed effects absorbing the familiar intraday pattern in spreads and volume. What identifies the coefficient is purely within-day, bucket-to-bucket variation; standard errors are clustered by stock and by date, where the dependence actually lives.

The coefficient on lagged-bucket clustering in the effective-spread regression is 0.0189, significant at the 1% level. Coefficients here are small by construction and should be read that

Table 14: Clustering and liquidity within the trading day

Outcome	Regressor	Coefficient	SE	Bucket-rows
Effective spread (log)	M_{b-1}^0	0.0189***	(0.0022)	612,672
Effective spread (log)	$M_{b-1}^{BLarge0}$	0.0242***	(0.0034)	252,848
Price impact, 60s (bp)	M_{b-1}^0	0.1041*	(0.0537)	612,670
Price impact, 60s (bp)	$M_{b-1}^{BLarge0}$	0.0655	(0.0745)	252,848

Thirty-minute buckets. Stock-day and bucket-of-day fixed effects, so all daily variation and the average intraday pattern are absorbed and only within-day, bucket-to-bucket variation identifies the coefficient. Standard errors clustered by stock and by date. Each specification includes the outcome's own lagged bucket and log turnover. * $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$.

way: this design deliberately throws away all the variation the daily panel uses. This is the version of the liquidity relationship that survives every re-specification tried in this study.

Order flow, conditional on clustering. The execution-relevant version of the impact result asks whether the *same* order-flow imbalance moves the midquote more on high-clustering days. The interaction term is 0.0514, not statistically distinguishable from zero. A positive value would mean the book is easier to push when round-price trading is heavy; the estimate here does not resolve the question in either direction.

Verdict. Of the paper's mechanism variables, the only one confirmed is the one that requires no inference: standing round-price depth. The inferred flow shares are nulls under heavy censoring, which is uninformative rather than damning. The within-day relationship stands (0.0189, $t = 8.80$) in a design where nothing slow-moving about a stock can produce it.

8 How fragile is the daily finding? Robustness

8.1 The placebo

This is the test that could have ended the study. The measure is a share of volume at a particular last digit. If a share of volume at *any* last digit predicts liquidity equally well, then nothing here is about round numbers — it is an artefact of how a concentration measure behaves, and the whole interpretation collapses.

Rebuilding the measure on each of digits one through nine and re-running the same regression, the round-price coefficient is 0.0188 ($t = 2.51$), and 0 of the nine placebo digits reach conventional significance. Digit five is flagged separately in Table 15 because half-way points are semi-focal and attract some of the same behaviour — treating it as a pure placebo would be stacking the deck.




Figure 5: The effective-spread regression re-run with the measure rebuilt on each last price digit in turn. The round-price digit is highlighted; bars are 95% confidence intervals.

Table 15: Placebo test: the same regression run on every last price digit

Last digit	Coefficient	SE	<i>t</i>
0	0.0188**	(0.0075)	2.51
1	-0.0125	(0.0092)	-1.36
2	0.0064	(0.0103)	0.63
3	0.0050	(0.0086)	0.58
4	-0.0020	(0.0098)	-0.20
5	-0.0017	(0.0091)	-0.19
6	-0.0008	(0.0097)	-0.09
7	-0.0091	(0.0094)	-0.97
8	-0.0125	(0.0090)	-1.40
9	-0.0092	(0.0098)	-0.94

Outcome is the log effective spread in the dynamic specification of Table 10, with the round-price share replaced by the share of buy-initiated large-trade volume at each last digit in turn. If the round-price result reflects stale limit orders at psychologically salient prices, digits other than zero should carry no comparable information. Digit five is marked as semi-focal because half-way points attract some of the same behaviour.

8.2 Subsamples, deeper dynamics, and an alternative estimator

Table 16: Robustness of the clustering–liquidity relationship

Specification	Coefficient	SE	Stock-days
Baseline (dynamic)	0.0188**	(0.0075)	74,034
Size Q1	0.0140	(0.0156)	15,062
Size Q2	0.0313**	(0.0127)	14,999
Size Q3	0.0217	(0.0146)	15,099
Size Q4	-0.0119	(0.0174)	14,743
Size Q5	0.0207	(0.0241)	14,054
Period H1 (2025-01 to 2025-03)	0.0164*	(0.0096)	36,842
Period H2 (2025-03 to 2025-04)	0.0037	(0.0103)	37,192
Excluding 2025-04-03 to 2025-04-11	0.0189**	(0.0079)	67,970
Five lags of the outcome	0.0124*	(0.0073)	70,014
0.1-yen-tick days only	-0.0033	(0.0613)	2,057

Each row re-estimates the dynamic specification of Table 10 with the log effective spread as the outcome, on the stated subsample or with the stated change. Size quintiles are formed once per stock on median market capitalisation. The period rows split the panel at its median trading date, derived from the sample rather than from a calendar. The event-window row removes the named week. The five-lag row replaces the single lag of the outcome with five, so the coefficient reads against a richer version of the outcome’s own history. * $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$.

Table 16 re-estimates the headline specification on subsamples and under two changes of specification. Across size quintiles the point estimates bracket the baseline and 1 of five reach significance on their own; at a fifth of the sample each, individual significance is not the expectation, and the informative fact is the absence of a clean size pattern in where the association lives.

Sub-periods. Neither half is individually significant (0.0164 and 0.0037); at half the sample each, that is expected under the pooled estimate as well as under a null, and the split is uninformative on its own. Removing 2025-04-03 to 2025-04-11 — the sample’s most violent week — leaves the estimate at 0.0189 (significant at the 5% level), so the crash days themselves are not the driver; whatever weakened the second half is spread across it.

Deeper dynamics. Against five lags of the outcome the coefficient falls to 0.0124 (significant at the 10% level). The claim “adds information beyond the outcome’s own history” therefore depends on history meaning one day; a richer dynamic absorbs much of it.

An estimator with different assumptions. A Fama–MacBeth cross-check (Fama & MacBeth, 1973) estimates the relationship separately in each daily cross-section and averages the estimates, with Newey–West standard errors (Newey & West, 1987) over the resulting series. It gives -0.0412 with $t = -6.44$ over 77 cross-sections — *opposite* in sign to the panel estimate, and

reliably so. The two designs answer different questions: the fixed-effects panel compares a stock with its own recent past, while the daily cross-sections compare stocks with each other. The positive spread association is a within-stock phenomenon; across stocks, conditional on yesterday’s spread, more-clustered names trade marginally tighter. A cross-check that disagrees is information about where the identification lives, not corroboration — any use of this indicator that ranks stocks against each other is using the sign that points the other way.

Verdict. The result is specific to the round-price digit rather than to the shape of the measure. The daily association itself is thinner than a pooled table suggests: the pooled estimate is significant at the 5% level, but it falls below conventional significance once the outcome’s history means five lags rather than one; its sign reverses between the within-stock and cross-sectional designs. The within-day relationship of Section 7 is the finding that survives every re-specification tried here, and conclusions are weighted accordingly.

9 A costed strategy demonstration

The brief behind this study is to build a rebalancing strategy on this indicator, and the distance between “the indicator predicts liquidity” and “here is a tradable rule” is exactly where such projects fail. So the machinery is built and run: form a signal from information available at the time, sort on it, rebalance, and charge realistic costs.

Table 17: Next-day returns to a size-neutralised clustering signal

Portfolio	Return (bp/day)	SE	<i>t</i>	Names	Eff. half-spread (bp)
Q1	+6.46	(17.56)	0.37	191	11.09
Q2	+3.58	(17.14)	0.21	192	8.73
Q3	+2.64	(17.22)	0.15	192	7.72
Q4	+2.29	(16.99)	0.14	192	8.07
Q5	+3.34	(16.91)	0.20	191	10.52
Q5 – Q1 gross, daily	-3.12	(2.16)	-1.44	–	10.80
net of costs	-35.09	–	–	–	–
Q5 – Q1 gross, 5-day signal	-2.03	(3.09)	-0.66	–	10.74
net of costs	-15.54	–	–	–	–

The signal is the daily cross-sectional residual of $M^{B\text{Large}0}$ on the same-day opening-digit contamination controls, relative tick size, log market capitalisation, lagged log turnover, lagged volatility, lagged spread and the overnight return, standardised within the day; the 5-day variant sorts on a five-day rolling mean of that residual. Portfolios are equally weighted and rebalanced daily. Costs charge the average measured effective half-spread of the two traded legs on both sides of every replacement in both legs, which is four times turnover times the legs’ half-spread. Stock-days with an absolute move above 25% are dropped, since closing prices here carry no adjustment for splits or dividends – note the March fiscal-year-end ex-dividend dates sit inside the sample. This is a demonstration that the pipeline runs end to end, not evidence of a profitable strategy: one short sample, one market, and no out-of-sample period.

The signal is the daily cross-sectional residual of the clustering measure on everything that mechanically drives it — the same-day opening-digit contamination, relative tick size, size,




Figure 6: Cumulative return to the long-short quintile portfolio, gross and net of estimated trading costs, for the daily and the five-day-smoothed signal.

turnover, volatility and spread — standardised within the day. Residualising is essential rather than cosmetic: clustering is strongest in small, retail-heavy names, so an unneutralised sort is largely a small-cap bet wearing a microstructure costume.

The cost line is the finding. The daily long-short portfolio replaces 74.0% of each leg per day. Every replacement is two trades and there are two legs, so the daily bill is four times turnover times the half-spread the traded legs actually quote (10.80 basis points — the extreme quintiles are wider than the panel average of 9.17): 31.97 basis points against a gross return of -3.12, leaving -35.09. A signal that has to be refreshed daily in names this wide pays its gross return away several times over.

Smoothing separates noise from signal. The daily signal's first-order autocorrelation is 0.11 — most of a day's reading is transitory — so the churn diagnosis is testable: sorting on a five-day rolling mean of the same residual cuts turnover from 74.0% to 31.5% of each leg per day and the cost line from 31.97 to 13.51 basis points, while the gross return goes from -3.12 to -2.03 ($t = -0.66$). Even at a fraction of the cost the net line does not turn positive, which is consistent with the daily reading being noise around an effect too small to trade outright.

That is the useful lesson for the problem this prototype was built for. If the indicator is to be used at all, it has to be used either at a horizon long enough to amortise the spread, or inside a rebalance that was going to happen anyway — which is precisely the execution-scheduling application the impact results in Section 5 point to, and precisely not a standalone alpha.

Verdict. The pipeline runs end to end and the economics are transparent. This is not evidence of a profitable strategy, and it is not evidence of an unprofitable edge either: four months, one market, no out-of-sample period, returns from closing prices with no adjustment for splits or

dividends (the March fiscal-year-end ex-dividend dates sit inside the window), which is why stock-days with moves beyond 25% are dropped rather than believed.

10 What to take to MTEC

10.1 What was found

Ohta's measures survive out of sample. On a year his study did not cover, round prices take 14.95% of large-trade volume against 10.77% of small-trade volume; the large-minus-small ordering his mechanism predicts holds (significant at the 1% level); finer grids cluster more; the paper's size gradient does not appear in this screened universe — documented with its tick-mix decomposition in Section 5 rather than assumed away; and trades at round prices carry a positive impact premium of 0.43 basis points at one minute — about 34% of the paper's published magnitude, and transient where his horizon ends. None of the comparisons was tuned; the benchmarks were fixed before the numbers were computed.

On the question the paper leaves open, clustering carries a measurable association with next-day trading costs beyond the outcome's own one-day lag (0.0188 on the log effective spread, $t = 2.51$; depth complements it at -0.0441), and a placebo across the other last digits leaves 0 of eight at conventional significance. The robustness work bounds the claim: the pooled estimate is significant at the 5% level, but it falls below conventional significance once the outcome's history means five lags rather than one; its sign reverses between the within-stock and cross-sectional designs. The within-day version — 0.0189 ($t = 8.80$) under stock-day and bucket fixed effects — is the form of the relationship this study would defend without qualification.

The book-based measures are the clearer contribution, read with discipline. Round prices hold 13.21% of standing visible depth against a uniform 10%, observing the stale inventory directly; standing round-price depth predicts round-price execution ($t = 11.4$), though partly mechanically; and the inferred flow shares are nulls under severe ten-level censoring, which adjudicates nothing. The inferred submission shares land on the paper's levels without calibration (0.102 and 0.103 against his 0.106 and 0.105), which is the strongest external validation in the study; the cancellation ratios overshoot (0.91 and 0.90 against 0.806 and 0.812), a real difference — more cancellation in 2025, or netting bias in the ten-level window — that is reported as open rather than absorbed into the validation claim.

The strategy demonstration prices the idea rather than selling it: gross -3.1 basis points a day against costs of 32.0 at daily rebalancing, and smoothing to five days cuts the costs to 13.5 without uncovering alpha. The economics point at execution, not selection.

10.2 What it is not

Not causal. Four months, no experiment, no instrument. Everything reported is a descriptive association or a predictive relationship. Clustering and liquidity are plausibly both driven by things this study does not observe, and nothing here identifies a direction of causation.

Not Ohta’s identification. He identifies noise traders with margin balances and ownership shares. Those are unavailable here, so this study leans entirely on the claim his conclusion advances — that the clustering measures are themselves the proxy. That makes the exercise a test of his proposal rather than an independent confirmation of it, and a genuine replication would need the ownership data.

A narrow window on the book. Ten visible levels sound like a lot and are not. On the stock-days examined in detail, 81% of quoted volume sits beyond them. The inferred order flow describes the neighbourhood of the best quote, which is where execution happens, but it is silent about the far book — and the far book is exactly where Ohta argues the interesting stale orders sit.

A sample shaped by its filters. Excluding non-power-of-ten tick sizes removes the 0.5-yen grid, and with it much of the mid-priced TOPIX 500 population; the 9:10 rule thins the panel on shock days. Both follow from the paper’s filters rather than from choices made here, but together they mean the regression sample is not the market, and is least representative exactly when markets are wildest.

10.3 Mapping the results onto the internship brief

The brief — a rebalancing strategy built on this indicator — has three defensible readings, and the results here speak to each differently.

(A) Execution-aware rebalancing. The best-supported use. The round-price impact premium is positive, front-loaded and transient (Section 5), and the within-day clustering–spread link is the study’s most robust regression (Section 7). Both are statements about *what an order pays and when*, which is what an execution schedule consumes: avoid crossing at round levels, avoid resting own orders at them, time non-urgent legs away from high-clustering buckets. No forecasting is required and a basis point is a real saving.

(B) A signal tilt. The least supported. The daily association is real but the pooled estimate is significant at the 5% level, but it falls below conventional significance once the outcome’s history means five lags rather than one; its sign reverses between the within-stock and cross-sectional designs, the cross-sectional sign points the other way (the Fama–MacBeth estimate), and the

costed demonstration in Section 9 shows the daily version pays its gross away several times over while the smoothed version has nothing left to pay with. If a tilt is wanted anyway, it needs months of holding period, size discipline, and an out-of-sample year.

(C) Rebalance-trigger design. Untested here and testable with this exact pipeline: modulate rebalancing bands or urgency by the residualised clustering state, on the argument that rebalancing is liquidity provision and the counterparty is more likely uninformed when noise traders are active. The panel and cost machinery this prototype builds is the spine that experiment needs.

Data to ask for on day one. Margin-trading balances (the *nisshoukin* securities-finance series) and ownership shares would convert the proxy test into a joint test of proxy and mechanism; a second year of tape would give the sub-period and out-of-sample splits this sample is too short for; and the full FLEX feed would lift the ten-level censoring that mutes the order-flow variables.

10.4 Conclusion

The sentence at the end of Ohta (2026) proposes that price clustering can serve as an observable proxy for noise-trader activity, and that its relationship to price formation and liquidity could then be studied at daily or shorter frequency. On the January–April 2025 Tokyo Stock Exchange tape the proposal survives its first contact with data the paper never saw: the measures replicate, the within-day liquidity relationship is robust, and the standing round-price inventory the mechanism requires is directly visible in the book. The daily version of the relationship is thinner — the pooled estimate is significant at the 5% level, but it falls below conventional significance once the outcome’s history means five lags rather than one; its sign reverses between the within-stock and cross-sectional designs — and the indicator does not support a standalone trading rule at this horizon, because the costs exceed the signal. The most defensible use of the measure is the one that needs no forecast at all: knowing what a round-price execution costs, and scheduling around it.

References

- Ahn, H.-J., Cai, J., & Cheung, Y. L. (2005). Price clustering on the limit-order book: Evidence from the stock exchange of Hong Kong. *Journal of Financial Markets*, 8(4), 421–451.
- Amihud, Y. (2002). Illiquidity and stock returns: Cross-section and time-series effects. *Journal of Financial Markets*, 5(1), 31–56.
- Aquilina, M., Budish, E., & O'Neill, P. (2022). Quantifying the high-frequency trading ‘arms race’. *The Quarterly Journal of Economics*, 137(1), 493–564.
- Barber, B. M., & Odean, T. (2008). All that glitters: The effect of attention and news on the buying behavior of individual and institutional investors. *The Review of Financial Studies*, 21(2), 785–818.
- Bhattacharya, U., Holden, C. W., & Jacobsen, S. (2012). Penny wise, dollar foolish: Buy–sell imbalances on and around round numbers. *Management Science*, 58(2), 413–431.
- Biais, B., Hillion, P., & Spatt, C. (1995). An empirical analysis of the limit order book and the order flow in the Paris Bourse. *The Journal of Finance*, 50(5), 1655–1689.
- Box, T., & Griffith, T. (2016). Price clustering asymmetries in limit order flows. *Financial Management*, 45(4), 1041–1066.
- Cameron, A. C., Gelbach, J. B., & Miller, D. L. (2011). Robust inference with multiway clustering. *Journal of Business & Economic Statistics*, 29(2), 238–249.
- Cao, C., Hansch, O., & Wang, X. (2009). The information content of an open limit-order book. *Journal of Futures Markets*, 29(1), 16–41.
- Chung, K. H., Van Ness, B. F., & Van Ness, R. A. (2004). Trading costs and quote clustering on the NYSE and NASDAQ after decimalization. *Journal of Financial Research*, 27(3), 309–328.
- Cont, R., Kukanov, A., & Stoikov, S. (2014). The price impact of order book events. *Journal of Financial Econometrics*, 12(1), 47–88.
- Davis, R. L., Van Ness, B. F., & Van Ness, R. A. (2014). Clustering of trade prices by high-frequency and non-high-frequency trading firms. *Financial Review*, 49(2), 421–433.
- Fama, E. F., & MacBeth, J. D. (1973). Risk, return, and equilibrium: Empirical tests. *Journal of Political Economy*, 81(3), 607–636.
- Foucault, T. (1999). Order flow composition and trading costs in a dynamic limit order market. *Journal of Financial Markets*, 2(2), 99–134.
- Glosten, L. R., & Milgrom, P. R. (1985). Bid, ask and transaction prices in a specialist market with heterogeneously informed traders. *Journal of Financial Economics*, 14(1), 71–100.
- Harris, L. (1991). Stock price clustering and discreteness. *The Review of Financial Studies*, 4(3), 389–415.
- Hasbrouck, J. (1991). Measuring the information content of stock trades. *The Journal of Finance*, 46(1), 179–207.

- Hendershott, T., Jones, C. M., & Menkveld, A. J. (2011). Does algorithmic trading improve liquidity? *The Journal of Finance*, 66(1), 1–33.
- Japan Exchange Group. (2024). *Tick sizes (Yobine no tan'i) — trading rules for domestic stocks*. Retrieved July 29, 2026, from <https://www.jpx.co.jp/equities/trading/domestic/07.html>
- Kaniel, R., Saar, G., & Titman, S. (2008). Individual investor trading and stock returns. *The Journal of Finance*, 63(1), 273–310.
- Kuo, W.-Y., Lin, T.-C., & Zhao, J. (2015). Cognitive limitation and investment performance: Evidence from limit order clustering. *The Review of Financial Studies*, 28(3), 838–875.
- Kyle, A. S. (1985). Continuous auctions and insider trading. *Econometrica*, 53(6), 1315–1335.
- Lee, C. M. C., & Ready, M. J. (1991). Inferring trade direction from intraday data. *The Journal of Finance*, 46(2), 733–746.
- Lehmann, B. N., & Modest, D. M. (1994). Trading and liquidity on the Tokyo Stock Exchange: A bird's eye view. *The Journal of Finance*, 49(3), 951–984.
- Næs, R., & Skjeltorp, J. A. (2006). Order book characteristics and the volume–volatility relation: Empirical evidence from a limit order market. *Journal of Financial Markets*, 9(4), 408–432.
- Newey, W. K., & West, K. D. (1987). A simple, positive semi-definite, heteroskedasticity and autocorrelation consistent covariance matrix. *Econometrica*, 55(3), 703–708.
- Nickell, S. (1981). Biases in dynamic models with fixed effects. *Econometrica*, 49(6), 1417–1426.
- Ohta, W. (2006). An analysis of intraday patterns in price clustering on the Tokyo Stock Exchange. *Journal of Banking & Finance*, 30(3), 1023–1039.
- Ohta, W. (2026). Noise traders and price clustering. *Gendai Finance*. <https://doi.org/10.24487/gendaifinance.480001>
- Peress, J., & Schmidt, D. (2020). Glued to the TV: Distracted noise traders and stock market liquidity. *The Journal of Finance*, 75(2), 1083–1133.
- Yamada, M. (2023). On high-frequency trading in the Japanese stock market. *TCER Working Paper Series*, (J-25).

Appendix

Appendix A Variable definitions

Notation follows Ohta (2026) where the quantity is his. A superscript $k \in \{B, S\}$ denotes trades initiated by a buy or a sell market order; *large* means a trade of more than one trading unit and *small* exactly one unit. All prices are handled as integer tenths of a yen, so the digit arithmetic is exact on every grid.

Clustering

<i>digit</i>	The last price digit, $(p \bmod 10\tau)/\tau$ for tick size τ . A <i>round price</i> is one whose digit is zero. Under a uniform distribution each digit would carry 10% of volume.
M^0	Share of continuous-session traded volume executing at a round price.
$M^{kLarge0}$	The same share within large trades initiated on side k . This is the paper's headline measure and the primary regressor here.
$M^{kSmall0}$	The same within small trades. The mechanism predicts it behaves differently, and it serves as the internal contrast.
M^d	Share of volume at last digit d , for $d = 0, \dots, 9$. Digits 1–9 are used only for the placebo test.
<i>Small</i>	Share of volume in trades of exactly one trading unit.

Liquidity, as the paper defines it

<i>ES</i>	Effective half-spread: the volume-weighted mean of $q(p - m)/m$ in basis points, where m is the midquote prevailing immediately before the trade and q is +1 for buy-initiated and -1 for sell-initiated trades.
$Imp(\Delta)$	Price impact: the volume-weighted mean of $q(m_{t+\Delta} - m_t)/m_t$ in basis points, at $\Delta \in \{1s, 60s, 300s\}$. Trades whose horizon would run past the end of their own session are excluded for that horizon, so no impact is ever measured across the lunch break or the close.
$RS(\Delta)$	Realized spread, $ES - Imp(\Delta)$, computed on that horizon's own trade set so the decomposition closes exactly.
ΔImp^{kLarge}	The round-price impact premium: impact of large round-price trades minus impact of other large trades, on side k . Cells with fewer than five trades are set missing.
Dep^{ask0}, Dep^{bid0}	Log mean depth at the best quote immediately before round-price trades, with

the non-round counterparts Dep^{ask1} , Dep^{bid1} .

QS	Time-weighted quoted spread in basis points, sampled on a one-minute grid over snapshots showing an ordinary two-sided book.
$Amihud$	Absolute open-to-close return divided by turnover in millions of yen, after Amihud (2002). Auxiliary only: it is included as a familiar yardstick, not as a substitute for the measures above.

Price formation

RV	Realised variance of five-minute midquote log returns, computed within each session so no return spans the lunch break or the overnight gap.
VR	Variance ratio, the variance of five-minute returns over five times the variance of one-minute returns. $ VR - 1 $ measures departure from a random walk in either direction.
OFI	Order-flow imbalance at the best quote, following Cont et al. (2014), accumulated over quote updates. Implemented from the published definition rather than taken from a library.

The visible ten-level book

$RDepth^{ask0}, RDepth^{bid0}$	Share of visible ten-level resting volume sitting at round prices, time-averaged on a one-minute grid. This is the measure Ohta's data could not support: it observes round-price orders while they are still standing, rather than after they have been executed against. Under a uniform digit distribution it would be 10%.
$Slope$	Depth slope: the regression coefficient of level volume on level index, averaged across the two sides, after Næs and Skjeltorp (2006).
L^{k0}	Round-price share of limit-order volume submitted on side k , inferred from ladder deltas within the visible book (Appendix B).
L^{k0C}	Cancelled volume at round prices divided by submitted volume at round prices; L^{k1C} is the same for other prices. The hypothesis is that round-price orders are cancelled <i>less</i> , because they are executed against first.
L^{k0E}	Executed volume at round prices divided by submitted volume there.
C^{k0}	Round-price share of all cancelled limit-order volume.
<i>distance classes</i>	Submissions are additionally split by where they sit relative to the prevailing best quote: inside it, at it, outside by less than 0.5%, and outside by 0.5% or more.

Controls and sample variables

τ	Tick size for the stock-day, taken from the exchange's published schedule (Japan Exchange Group, 2024) as encoded in this repository, and cross-checked against the finest increment observed on the tape.
<i>RelTick</i>	Tick size divided by the opening price, in basis points.
$D^n \times Lowvol$	Interaction of a dummy for the opening price's last digit being n with a dummy for realised variance falling below that day's cross-sectional median, controlling the mechanical contamination documented in Ohta's Table 2 and reproduced in Section 5. The timing matches the regressor being cleaned: specifications using M_{t-1} carry $D^n_{t-1} \times Lowvol_{t-1}$, the contemporaneous specification and the strategy signal carry the same-day interaction. Digit zero is the omitted category.
ret^{on}	Overnight return, previous close to opening price.
<i>Turnover</i>	Traded value in yen; entered in logs.
<i>MktCap</i>	Closing price times shares outstanding, from the daily summary product; entered in logs and used to form size quintiles.
<i>FineTick</i>	Indicator for a stock-day on the 0.1-yen grid. These days enter the descriptive results and are excluded from the regressions, following the paper.

Appendix B Inferring limit-order flow from the visible ladder

The feed publishes the ten best price levels per side after every event, but not the orders themselves. Submission and cancellation volume must therefore be inferred from how the displayed depth changes. This appendix states the accounting, the two rules that make it correct, what it cannot see, and the test that checks it.

B.1 The accounting identity

Between two consecutive snapshots, the volume resting at a given price can change only by execution, cancellation, or new submission:

$$V_n(p) = V_{n-1}(p) - \text{executed}_n(p) - \text{cancelled}_n(p) + \text{submitted}_n(p). \quad (9)$$

Writing $\delta = V_n(p) - V_{n-1}(p)$ and letting $X_n(p)$ be the volume printed at p in that interval, equation (9) rearranges to

$$\text{submitted}_n(p) = \max\{\delta + X_n(p), 0\}, \quad \text{cancelled}_n(p) = \max\{-(\delta + X_n(p)), 0\}. \quad (10)$$

Netting within the interval means an interval containing both a submission and a cancellation at the same price reports only the difference. This biases both flows downward, symmetrically, and the share statistics built from them are therefore more robust than their levels.

B.2 Two rules that make it correct

Match by price, never by level index. When the book shifts a level — the best ask is consumed and everything moves up one slot — an index-based comparison sees every level change and manufactures a submission and a cancellation at each. Matching by price sees what actually happened: one price emptied. This is the single largest source of error in ladder-based order-flow inference, and the test in Section B.4 exercises it directly.

Count a price only where it is observable in both snapshots. On the ask side, every price at or below the deepest visible ask is observable: if it is listed, its volume is shown, and if it is not listed, it is genuinely empty, since a populated price inside the window would have displaced one of the ten. Above that frontier nothing can be said. The comparison is therefore restricted to

$$p \leq \min\{P_{n-1}^{10}, P_n^{10}\}$$

on the ask side, and the mirror condition on the bid side. Without this rule, a price entering the window from beyond level ten arrives carrying pre-existing depth that would be recorded as an enormous submission.

Two further restrictions are applied for the same reason. An interval is used only when both of its endpoints show an ordinary two-sided quote, so nothing is inferred across a special quote or a halt, when the displayed book does not mean what it usually means; and only when the two rows are adjacent on the tape and in the same session, so that a dropped stretch is never bridged. One unusable snapshot therefore invalidates the two intervals that touch it. That is deliberately conservative: it discards real flow rather than inventing any.

B.3 What this cannot see

Three limits are worth stating plainly, because they bound how the results in Section 7 should be read.

1. **Depth beyond ten levels is invisible.** On the stock-days examined in detail, between 95% and 98% of quoted volume sits outside the ten displayed levels, in the aggregated OVER and UNDER buckets. Ten levels of a 1-yen grid around a 4,000-yen price span a quarter of one percent, so this is arithmetic rather than a defect — but it means the inferred flow describes the neighbourhood of the best quote, not the whole book. This matters most for the furthest distance class, which is precisely where Ohta argues the interesting orders sit, and it is the main reason the measures here are treated as complements to his rather than replacements.
2. **Order modifications are indistinguishable from a cancel-and-submit pair** at two different prices. For the digit shares used here this is harmless, since both the cancellation and the submission are attributed to the prices they actually occurred at.
3. **Off-book trading is excluded entirely.** Only the central limit order book is observed.

One point cuts the other way and is worth recording: the TSE central book has no hidden or iceberg orders, so displayed depth is the whole of the resting depth at each visible price. The inference is exact where it applies.

B.4 The test

Correctness is checked against a hand-built order book whose flow is known exactly, rather than against another implementation. The tape is six snapshots long on a 1-yen grid with the ask side starting at 1,000, 1,001 and 1,002, and it is constructed so that each of the three known failure modes would produce a visibly wrong answer:

1. 300 shares are added at 1,000 — a submission, at the round price.
2. A buyer lifts 500 shares at 1,000. Displayed depth falls by 500, but the print explains all of it, so neither a submission nor a cancellation may be inferred. An implementation that forgets to net the execution reports 500 shares of phantom cancellation.
3. 200 shares vanish at 1,000 with no print — a cancellation.
4. The remaining 100 shares leave, 1,000 empties, the ladder shifts up, and 1,003 enters from beyond the old window. The 100 must be counted; the depth arriving at 1,003 must not,

and index matching would additionally report spurious flow at all three levels.

The test asserts the exact expected totals — 300 submitted, 300 cancelled, both entirely at the round price, an execution ratio of 500/300, and no flow at all on the untouched bid side — together with a separate case confirming that a special quote in the middle of the tape invalidates the two intervals touching it rather than being bridged across.

B.5 Does it recover the published magnitudes?

The strongest available check is external, and it is the one this appendix rests on. Ohta reports a mean L^{S0} of 0.106 and a mean L^{S0C} of 0.806 across 2010–2022, computed from the same kind of feed by an independent implementation.

Run over the panel, the inference returns a mean round-price share of submitted sell limit orders of 0.1024 and of buy limit orders 0.1025, against Ohta’s published 0.106 and 0.105 — agreement to half a percentage point, from a different sample period, with nothing calibrated. The cancellation ratios are 0.909 and 0.905 against his 0.806 and 0.812: a genuine overshoot, consistent with either faster cancellation in 2025 or with netting bias inside the ten-level window, and treated as an open difference rather than folded into the validation claim. These figures come from 9,855 stock-days.

Agreement on the submission-share levels is not something a broken implementation produces by accident: the two rules in Section B — matching by price and requiring observability in both snapshots — are exactly what stands between this and a number several times too large. The cancellation-ratio overshoot is the honest residual of the exercise. It has two live explanations — cancellation traffic genuinely rose between the paper’s sample and this one, or the netting inside ten-level intervals biases the ratio upward here — and nothing in this study distinguishes them, so the ratio is not claimed as validated.

Appendix C Reproducibility

C.1 Pipeline

The study runs as eight numbered stages. Each begins with a gate that refuses to proceed unless the previous stage's outputs are present and correct, and ends by writing a report of what it did. Long stages checkpoint per unit of work, so an interrupted run resumes rather than restarting.

- S0** Environment, trading calendar, and the exchange tick-size schedule. The panel estimators are verified against a simulated panel with a known coefficient before any real data is fitted.
- S1** The measure library, validated against the anchor stock-days and the hand-built order book.
- S2** Universe selection, then ingest of the raw feed into a columnar store.
- S3** The panel build: one read per stock-day produces the daily row, the thirty-minute buckets, and the ladder aggregates.
- S4** Replication of the paper's stylized facts. This stage carries a stop rule.
- S5** Panel regressions and robustness.
- S6** The strategy demonstration.
- S7** This report.

C.2 The stop rule

Three of the paper's findings are treated as necessary conditions rather than results: clustering must be higher in large trades than small ones, the round-price impact premium must be positive, and clustering must be elevated on the finest tick grid. These are not delicate effects. If a pipeline fails to reproduce them, the explanation is overwhelmingly likely to be a defect in the pipeline — an inverted trade-direction map, a floating-point digit, a mis-assigned tick — and not a change in the market. S4 therefore halts rather than proceeding to results that would be built on a broken measurement.

C.3 Tests

The suite covers the exchange tick-size bands and their boundary conventions (each case a tick observed on the tape rather than an invented one); the integer digit arithmetic, including a test that records how the naive floating-point version fails; the session-close change of 2024-11-05; the trade-signing map including the stop-quote variants that decode to an unrecognised label; the write guard that keeps the pipeline out of the read-only source tree; the tick-resolution rule, with the three real stock-days that motivated it; and the ladder algorithm against a hand-built book.

Two tests exist because their failure modes shipped once and were caught in audit rather than by the pipeline. The low-volatility indicator once degenerated to all-false through a

window-scoping accident, which silently zeroed every opening-digit contamination dummy; the regression frame now refuses degenerate dummies, and a test pins the fix. And the sub-period robustness split was once hardcoded to a calendar date from an earlier study window, so that one half was empty and the other was the whole sample; the split is now derived from the panel's own dates, asserts that the halves partition the sample, and is tested on synthetic panels.

C.4 Two decisions worth recording

The tape outranks the schedule. The exchange's published tick-size table is a function of price and index membership, and point-in-time membership lists drift. Where the tape shows a finer grid than the table predicts, the tape is right by construction — a stock cannot quote inside a grid it is not on. Where the tape shows a *coarser* grid the case is ambiguous, because an illiquid stock skips grid points, so the tape is believed only when the stock-day supplied enough distinct prices for the observation to mean something. Getting this wrong is not a rounding error: three stock-days in an early test panel were forced onto a 1-yen grid while actually trading on 5- and 10-yen grids, and their clustering measures read 50% and 100% because every observed price then lands on a multiple of the assumed tick.

Ingest had to be restructured. The reader library's default ingest assembles a whole trading day in memory before writing. At current message rates a single day of the full market exceeds the memory of the machine used here, and the ingest fails. The library's streaming path — which writes each part of a day as it is read and then discards it — engages for requests that name a bounded set of stocks, so the universe is ingested in code-ordered batches at that threshold. Because the feed writes each day's parts in ascending stock-code order and the library prunes a filtered request to the parts spanning its codes, a code-ordered batch reads roughly its own slice of the tape rather than all of it.

One false start is worth recording because it is instructive. Raising the library's streaming threshold in the parent process appeared to work and made matters worse: ingest workers are spawned processes that re-import the library, so they never saw the raised value and continued assembling whole days, while the parent — which *had* seen it — sized their memory as though they were streaming and started more of them than would fit. The configuration that works is the one that leaves the threshold alone and keeps each request inside it.

C.5 Data

Source is the Nikkei NEEDS TICST120 product — tick data with the ten best quotes per side — read strictly read-only, together with the TICSS110 daily summaries for trading units, shares outstanding and the universe screen. No raw or derived data is included in the code repository or leaves the machine on which it was processed; the licence is institutional. The figures and aggregates in this report are the only data products that travel.